

# COMP 512 Project Deliverable 2

## Performance Evaluation

### Testbed description

#### Client

The client consists of series of basic transaction operations (i.e. a sequence that creates a flight, queries it, and reserves it). The client also has a main method that runs experiments by invoking the various transaction elements in different ways. For example it computes the average transaction time for  $n$  transactions with a load interval of  $m$  milliseconds. The client can also take a delay parameter which is used in the multiple clients experiment (figure 1).

#### Measurement techniques

The time for transactions is measured by taking the system time before the transaction is invoked and the system time after and using the difference.

To ensure that size of data set is not an issue, the data on the RMs was reset for every data point (e.g. for each transactions per second value, the whole system was reset).

#### Transactions used

We have two primary transaction sequences, one for a single RM test and one for a multiple RM test.

The single RM test is the following transaction:

NEWFLIGHT, QUERYFLIGHT, RESERVEFLIGHT, COMMIT

The multiple RM test is the following transaction:

NEWFLIGHT, NEWCAR, NEWROOM, COMMIT

### Experiments conducted

We conducted two experiments.

Experiment 1 (figure 1), Increasing load (i.e. transactions per second to middleware) across multiple clients (10 clients) and measuring response time for both single and multiple RMs

Experiment 2 (figure 2), Increasing number of consecutive iterations of a transaction submitted to middleware from a single client for both single and multiple RMs

### Performance figures

The performance figures begin on the following page.




Figure 1 – Multiple Clients

### Observations (Experiment 1)

Both Multiple and Single RMs start out relatively flat, then increases quickly before plateauing and changing at a slower, although more volatile pace. Multiple RM always has a slower average response time than Single RM for the same number of transactions submitted to the middleware per second.

### Analysis (Experiment 1)

- The higher response time for multiple RM is likely explained by the need to start “sub-transactions” on each RM and consequently, the higher number of IO operations (each RM start transaction / commit operation involves an IO operation on disk).
- The rapid growth is likely occurring around the point at which the delay between transactions (for controlling load) is reduced to 0 and operations are now reaching each of the RMs concurrently, also potentially triggering a competition for resources as only one operation may be sent from the client at a time due to the limits of Java Sockets.
- The rapid growth occurs earlier for Multiple RM because the requests take longer to serve on average and the concurrency on the RMs arrives earlier.
- The volatility after the rapid increase is likely the result of non-deterministic behaviour when transactions / operations become concurrent and must compete for resources. For example, a transaction may simply be unlucky and starts before other transactions but ends up finishing after the other transactions, dragging up the average response time.




Figure 2 – Single Client

### Observations (Experiment 2)

Both Multiple and Single RMs are somewhat consistent, then increases at a faster pace, although there is volatility in both cases, Multiple RM appears to be more volatile. Multiple RM has a higher response time than Single RM for the same number of iterations.

### Analysis (Experiment 2)

- The higher response time for multiple RM is likely explained by the need to start “sub-transactions” on each RM and consequently, the higher number of IO operations (each RM start transaction / commit operation involves an IO operation on disk).
- The appearance of consistency near the beginning may be noise. If so then the general trend upwards could be explained by increasing load on disk resources. As the number of iterations grow, so does the total size of the serialized file being written to disk (this is simply because the transaction was set up in such a way that each additional iteration of the transaction would add more data to the dataset). It is also possible other hardware limitations are at play here.
- The gap between single RM and multiple RM remains fairly constant which would suggest the cost of accessing multiple RMs scales with the transaction and not with the number of consecutive iterations.

## Conclusions

### Potential bottlenecks

Given the design of our system, it is likely the largest bottleneck is the IO operations involved in starting, aborting, and committing transactions.

Certainly, accessing multiple RMs incurs a penalty. However, it is unlikely this penalty is the result of network constraints as the whole system was evaluated on a single machine and each operation is already sent as a separate message. It is also unlikely that this is the result of CPU constraints because the multiple RMs are all running in the foreground anyway. So the only remaining difference between a transaction that involves one RM and a transaction that involves many RMs is the number of places the middleware must start (or enlist an RM for) a transaction and similarly how many places the RM must commit or abort. The commit mechanism is what likely ends up being the most costly because it requires that expensive disk operation.

Although it is difficult to speculate on the behaviour of the hardware (there may be optimizations done by the OS that I do not know about), figure 1 suggests that concurrency is not a problem. Namely that the Multiple RM time is about double the Single RM time both at the stable early part (low load) and also double after the rapid increase has occurred for both.

### Where is the time spent

It is likely that the time spent in the middleware is larger than that of the RM servers. This is because there is overhead on top of the middleware's own RM for customers and coordinator TM.

We see this in figure 2 where the multiple RM that involves 3 RMs and 1 middleware versus the single RM that involves 1 RM and 1 middleware. Despite doubling the number of transaction managers and I/O operations (the middleware also writes to disk on commit), the response time increases by less than double. Indeed, the gap between Single and Multiple remains fairly constant suggesting at least some portion of the time is not spent in the RMs.